



huggingfaceR
API-FIRST ML FOR R

merged to main

FOR R USERS

Hugging Face models, directly from R

No Python by default. Tidy inputs. Tidy outputs.
Examples you can run.

```
library(huggingfaceR)
hf_classify("I love using R")

# A tibble: 1 x 3
text      label    score
"I love..." POSITIVE 1.000
```

R

NATIVE WORKFLOWS

Use models from the Hub
without leaving your R
session.

chat

embeddings

structured extraction

audio

images

Hub APIs



huggingfaceR
TARGET USERS

R-first

WHO IT HELPS

For people who live in R and work with messy data

Text, documents, interviews, images, tickets, reviews,
notes, and research data.

Data scientists

Prototype model-backed workflows
without switching stacks.

01

Researchers

Code interviews, open-ended
survey answers, and document
collections.

02

Analysts

Classify, search, summarize, and
extract fields into tibbles.

03

Journalists

Transcribe clips, search archives,
and tag visual material.

04



● WHAT YOU CAN DO

Turn unstructured inputs into analysis- ready data

01 Classify text

Survey responses, reviews, support tickets, open-ended comments.

02 Create embeddings

Search, clustering, matching, topic discovery, nearest neighbors.

03 Extract fields

Pull names, themes, locations, sentiment, and codes from prose.

04 Analyze media

Transcribe audio, classify images, detect objects, generate images.



huggingfaceR
NEW UPDATE

feature expansion

WHAT CHANGED

A wider Hugging Face API surface, from one R package



Chat, streaming, tools

Use open-source chat models, stream responses, and connect model tool calls to R functions.



Structured output

Ask models for JSON-like fields and get one tidy row per input.



Multimodal wrappers

Audio transcription, text-to-image, image classification, and object detection.



Hub workflows

Download files, inspect provider options, search Spaces and papers, and guard writes.



huggingfaceR

DOCS NOW RENDER OUTPUT

● tested live

● NOT JUST SNIPPETS

The README and site now show **real output**

MLK audio

Transcription example

"I have a dream..." renders from a public Hugging Face audio sample.

COCO cats

Vision example

Image classification returns "Egyptian cat"; object detection draws cat boxes.

red cube

Image generation example

The README includes a generated image and the file output from R.



huggingfaceR
TRY IT

open source

WHERE TO START

What would you try first in R?

Survey text, interviews, embeddings, image tagging, or Hub downloads?

`farach.github.io/huggingfaceR`

`github.com/farach/huggingfaceR`

Pick a workflow. Run the example. Keep the result in a tibble.

01

Set token

02

Call model

03

Get tibble

04

Analyze in R